

# Final Review

courtesy of Joe Hellerstein for some slides

---

Jianlin Feng

School of Software

SUN YAT-SEN UNIVERSITY

---

# Topics Covered

- Relational Model
  - Disks, Files, Buffers, Indexes
  - Languages: Algebra, SQL
  - Query Processing: Sorting, Hashing, Join Algorithms
  - Query Optimization
  - Schema Refinement and Normalization
  - Concurrency Control
  - Crash Recovery
-

---

# Overview

- Purpose of this course: give students both
    - An understanding of what systems do, why we use them, and how to use common databases efficiently,
    - and an understanding of how databases work internally.
-

---

# Introduction

- What are databases?
  - Data models
  - What does a DBMS provide that the OS does not?
    - Levels of Abstraction, Data Independence, Concurrency Control, Crash Recovery, etc.
-

---

# Introduction (Contd.)

- schemas & data independence
    - conceptual schema
    - physical schema
    - external schema (view)
    - logical & physical data independence
-

---

# The Relational Data Model Basics

- Components of the model:
    - Relations, Attributes, Tuples
  - SQL Data Definition Language
  - Integrity Constraints
    - how do they come into being?
    - understand what you can learn from schema vs. instance!
  - *Referential Integrity*
    - a state which holds when all foreign key constraints are enforced.
-

# Relational Model (Contd.)

- Keys, Primary Keys, Foreign Keys, Candidate Keys
- Foreign key: Set of fields in one relation that is used to 'refer' to a tuple in another relation.
  - Must correspond to primary key of the second relation.
  - Like a 'logical pointer'.
- A set of fields is a key for a relation if :
  1. No two tuples can have same values in all key fields, and
  2. This is not true for any subset of the key.
  - Part 2 false? A superkey.
  - If there's >1 key for a relation, all are candidate keys. One of the candidate keys is chosen to be primary key.
- E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

# Memory Management

- Hierarchy of storage: RAM, Disk, Tape
- Advantages/disadvantages of different types of storage .
- Buffer management
  - You should know basics such as
    - Understand data structures required
    - notions of *replacement*, *dirty* pages, *pinning* pages
  - Understand different replacement policies
    - LRU, MRU, CLOCK
    - be able to simulate each, understand pros, cons!



---

# Memory Management (Contd.)

- Organizing records in pages
    - Fixed & variable-length fields in tuples
      - 2 alternatives for variable-length fields
    - Fixed & variable-length tuples on pages
      - know what a RID is, how it interacts with page layout
      - know details of “slotted page” with slot directory
  - Organizing pages in files
-

# File Organization

- Different file organizations: heap files, sorted files, hashed files
  - be able to compute costs of operations over each!
- Access costs for different organizations

# File Organization: Indexes

- Understand search keys (vs. key constraints!)

- 3 alternatives for data entries

Data record with key value  $k$

$\langle k, \text{rid of data record with search key value } k \rangle$

$\langle k, \text{list of rids of data records with search key } k \rangle$

- Choice of alternative orthogonal to indexing technique!


---

# Tree Structured Indexes

- Trees do *range* and *equality* search
  - ISAM, rules for adding and removing entries
  - B-Trees, rules for adding entries
    - You do not have to do “coalescing” on B+-tree deletion
-

# B+ Tree Structure

- Each node contains  $d \leq m \leq 2d$  entries (index or data)
  - The parameter  $d$  is called the *order* of the tree.
  - Each internal node contains  $m$  index entries: <key, page id>.
  - Each leaf node contains  $m$  data entries: <key, record *or* record id>
- The ROOT node contains between 1 and  $2d$  index entries.
  - It is a leaf or has at least two children.
- Each path from the ROOT to any leaf has the *same length*.
- Supports equality and range-searches efficiently.



---

# External Sorting, Hashing

- How to sort any file using 3 memory Pages
  - How to sort/hash in as few passes given some amount of memory
  - Relationship between buffers of memory, size of file, and number of passes to merge
  - Duality of sort and hash
  - Application to duplicate elimination, group by
-

# Relational Algebra

- Query language operating on relations
- Know the operators!
  - $\sigma$ ,  $\pi$ ,  $\times$ ,  $\cup$ ,  $\cap$ ,  $-$ ,  $\rho$ ,
  - know schemas of output relations
  - know varieties of joins
    - (conditional vs. equi vs. natural), division
  - be able to express complex operations in terms of simple ones
- Use relational algebra to express queries written in English, and vice versa.

# SQL

- DDL: “Create Table”
- DML: Delete From, Insert Into, Update
- Basic Query:  
select <targets>  
from <relations>  
where <qualification>
- Use of Distinct clause
- Set operations: Union, Except, Intersect

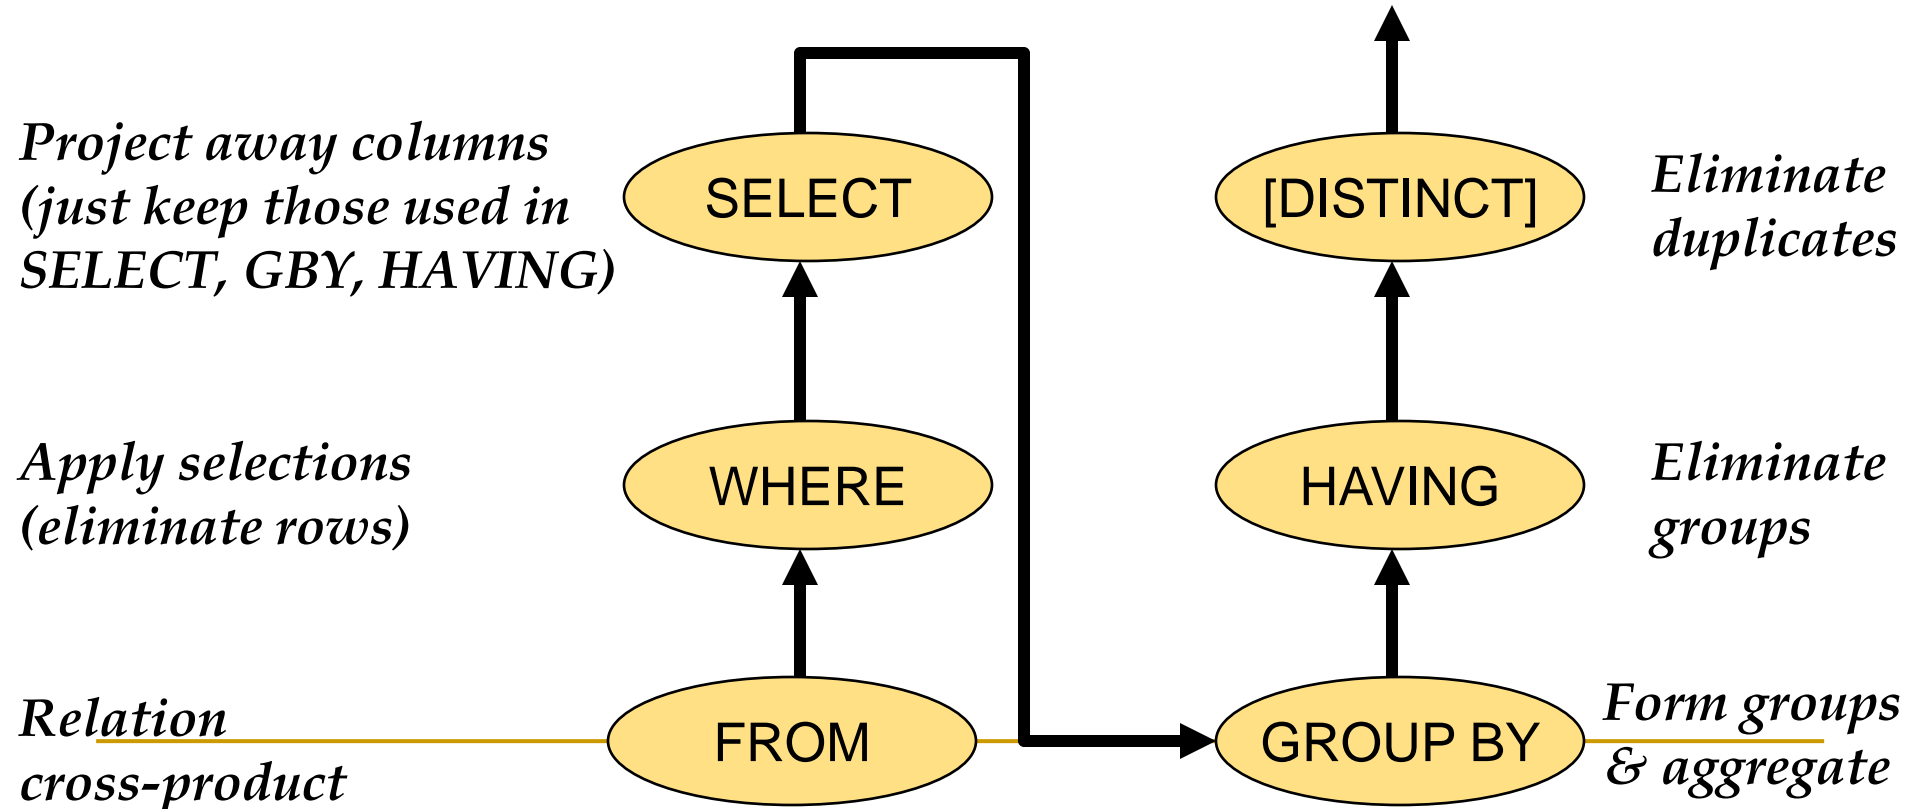


# SQL Query Language

- basic queries
- conceptual evaluation
- range variables
- expressions
- strings
- Union, Intersect, In, Except, Exists
- nested queries, correlated and not
- set comparison
- Aggregation
  - operators: Count, Avg, Any
  - Group By, Having

# Conceptual SQL Evaluation

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>



# Conceptual SQL Evaluation

SELECT	[DISTINCT]	target-list
FROM		relation-list
WHERE		qualification
GROUP BY		grouping-list
HAVING		grouping-qualification

**ORDER BY**

*Project away columns  
(just keep those used in  
SELECT, GBY, HAVING)*

SELECT

*Eliminate  
duplicates*

[DISTINCT]

*Apply selections  
(eliminate rows)*

WHERE

*Eliminate  
groups*


HAVING

*Relation  
cross-product*

FROM

*Form groups  
& aggregate*

GROUP BY



# SQL Continued

- Aggregation
  - Count, Sum, Avg, Max, Min
  - Group By, Having clauses
- Nested Queries
  - in Where or From Clause
  - set comparison:
    - In, Exists, Unique
    - *op* Any, *op* All
  - correlated vs. uncorrelated

# Implementation of Relational Operators

- Important Operators:
  - Selection
  - Projection
  - Set operations
    - Set-Difference, Union
  - Aggregation
  - Join
- Understand cost estimation, selectivity

# Selection

- with no index, scan entire relation
- with clustered index, use index
- with unclustered index, sort rids
- hash index only good for equality selection
- with multiple selection conditions, either
  - scan entire table testing all conditions
  - use index on most restrictive condition first, scan result for other conditions
  - use index for each condition, do set intersection on RIDs

---

# Projection

- Hard part: removing duplicates (if necessary)
  - Can remove duplicates by sorting or hashing
  - If index contains projected attr, can do index-only scan
-

# Set Operations & Aggregation

## ■ Set operations:

- ❑ Intersection, cross product treated like joins
- ❑ Union (Distinct) and Set Difference both involve finding duplicates between two sets, treat similar to Project

## ■ Aggregation

- ❑ without Group By, must scan entire relation
- ❑ with Group By, could sort, then scan
  - Or...?



# Joins

- Two relations: inner N and outer M
- Simple Nested Loops:
  - for each outer tuple, scan inner for matches
  - cost:  $M + \# \text{tuples in } M * N$
- Paged Nested Loops
  - for each page in M, scan inner for matches to any tuple in that M page
  - cost:  $M + M * N$
- Blocked Nested Loops
  - like paged, except put as much of M in memory as possible, leaving 1 page for N and 1 page for output
  - cost:  $M + \frac{M}{(\text{block size})} * N$

# Joins (Contd.)

## ■ Indexed Nested Loops

- for each tuple in M, use index to find matches in N
- cost:  $M + \text{\#tuples in } M * \text{cost to use index to get tuples}$

## ■ Sort-Merge Join

- Sort each table, merge finding like values,
- can be bad if many duplicates
- cost:  $M \log M + N \log N + M + N$

## ■ Hash Join

- partition both relations into buckets, read in one bucket from M at a time, match with elements from same bucket in N
- cost: partitioning  $2*(M + N)$  plus matching  $(M + N)$

# Query Optimization

- Some operations are commutative, associative
- May change the order of many operations in query
- Dramatic changes in cost depending on operation order
- “Query Plan” - a tree of operations indicating operation implementations and order
- Ideally find optimal plan
- In reality avoid terrible plans

# Optimizer Implementation

- need iterator interface so each operation passes tuples on to the next
- need cost estimator to determine costs of different plans
  - Reduction Factor of operations
- need statistics and catalogs to estimate costs

---

# System R Optimizer in Action

- convert SQL to relational algebra
  - find alternate plans
    - for each relation, consider all access paths
    - for multiple relations, consider different join algorithms, access paths
    - for join orders, only consider left-deep trees
-

# Functional Dependencies (FDs)

- An FD  $X \rightarrow Y$  holds over relation schema R if, for *every allowable instance*  $r$  of  $R$ :

$$t1 \in r, t2 \in r, \pi_X(t1) = \pi_X(t2)$$

$$\text{implies } \pi_Y(t1) = \pi_Y(t2)$$

( $t1, t2$  are tuples;  $X, Y$  are *sets* of attributes)

- In other words:  $X \rightarrow Y$  means
  - Given any two tuples in  $r$ ,  
if the  $X$  values are the same,  
then the  $Y$  values must be the same.  
(but not vice versa)

- Read “ $\rightarrow$ ” as “*determines*”

# Problems Due to $R \rightarrow W$

S	N	L	R	W	H
123-22-3666	Attishoo	48	8	10	40
231-31-5368	Smiley	22	8	10	30
131-24-3650	Smethurst	35	5	7	30
434-26-3751	Guldu	35	5	7	32
612-67-4134	Madayan	35	8	10	40

Hourly\_Emps

- Update anomaly: Can we modify W in only the 1st tuple of SNLRWH?
- Insertion anomaly: What if we want to insert an employee and don't know the hourly wage for his or her rating? (or we get it wrong?)
- Deletion anomaly: If we delete all employees with rating 5, we lose the information about the wage for rating 5!

# Rules of Inference

- **Armstrong's Axioms** ( $X, Y, Z$  are sets of attributes):
  - Reflexivity: If  $X \supseteq Y$ , then  $X \rightarrow Y$
  - Augmentation: If  $X \rightarrow Y$ , then  $XZ \rightarrow YZ$  for any  $Z$
  - Transitivity: If  $X \rightarrow Y$  and  $Y \rightarrow Z$ , then  $X \rightarrow Z$
- **Sound** and **complete** inference rules for FDs!
  - using AA you get *only* the FDs in  $F^+$  and *all* these FDs.
- Some additional rules (that follow from AA):
  - Union: If  $X \rightarrow Y$  and  $X \rightarrow Z$ , then  $X \rightarrow YZ$
  - Decomposition: If  $X \rightarrow YZ$ , then  $X \rightarrow Y$  and  $X \rightarrow Z$



# Attribute Closure

- Computing closure  $F^+$  of a set of FDs  $F$  is hard:
  - exponential in # attrs!
- Typically, just check if  $X \rightarrow Y$  is in  $F^+$ . Efficient!
  - Compute attribute closure of  $X$  (denoted  $X^+$ ) wrt  $F$ .  
 $X^+ =$  Set of all attributes  $A$  such that  $X \rightarrow A$  is in  $F^+$ 
    - $X^+ := X$
    - Repeat until no change:  
if  $U \rightarrow V \subseteq F, U \subseteq X^+$ , then add  $V$  to  $X^+$
  - Check if  $Y$  is in  $X^+$
  - Approach can also be used to find the keys of a relation.
    - If  $X^+ = R$ , then  $X$  is a superkey for  $R$ .
    - Q: How to check if  $X$  is a “candidate key”?

# Boyce-Codd Normal Form (BCNF)

- Reln R with FDs F is in BCNF if, for all  $X \rightarrow A$  in  $F^+$ 
  - $A \subseteq X$  (called a trivial FD), or
  - X is a superkey for R.
- In other words: “R is in BCNF if the only non-trivial FDs over R are key constraints.”

# Decomposition of a Relation Scheme

- How to normalize a relation?
  - *decompose* into multiple normalized relations
- Suppose  $R$  contains attributes  $A_1 \dots A_n$ .  
A decomposition of  $R$  consists of replacing  $R$  by two or more relations such that:
  - Each new relation scheme contains a subset of the attributes of  $R$ , and
  - Every attribute of  $R$  appears as an attribute of at least one of the new relations.

# Lossless Join Decompositions


- Decomposition of R into X and Y is lossless-join w.r.t. a set of FDs F if, for every instance  $r$  that satisfies F:

$$\pi_X(r) \bowtie \pi_Y(r) = r$$

- It is always true that  $r \subseteq \pi_X(r) \bowtie \pi_Y(r)$ 
  - In general, the other direction does not hold!
  - If it does, the decomposition is lossless-join.
- Definition extended to decomposition into 3 or more relations in a straightforward way.
- *It is essential that all decompositions used to deal with redundancy be lossless! (Avoids Problem #1)*

# Lossy Decomposition (example)

A	B	C
1	2	3
4	5	6
7	2	8




A	B
1	2
4	5
7	2

B	C
2	3
5	6
2	8

$A \rightarrow B; C \rightarrow B$

A	B
1	2
4	5
7	2




B	C
2	3
5	6
2	8

=

A	B	C
1	2	3
4	5	6
7	2	8
1	2	8
7	2	3

# Lossless Decomposition (example)

A	B	C
1	2	3
4	5	6
7	2	8




A	C
1	3
4	6
7	8

B	C
2	3
5	6
2	8

$A \rightarrow B$ ;  $C \rightarrow B$

A	C
1	3
4	6
7	8



B	C
2	3
5	6
2	8

=

A	B	C
1	2	3
4	5	6
7	2	8

But, now we can't check  $A \rightarrow B$  without doing a join!

## Dependency Preserving Decompositions (Contd.)

- Decomposition of  $R$  into  $X$  and  $Y$  is dependency preserving if  $(F_X \cup F_Y)^+ = F^+$ 
  - i.e., if we consider only dependencies in the closure  $F^+$  that can be checked in  $X$  without considering  $Y$ , and in  $Y$  without considering  $X$ , these imply all dependencies in  $F^+$ .
  - (just the formalism of our intuition above)
- Important to consider  $F^+$  in this definition:
  - $ABC, A \rightarrow B, B \rightarrow C, C \rightarrow A$ , decomposed into  $AB$  and  $BC$ .
  - Is this dependency preserving? Is  $C \rightarrow A$  preserved????
- Note:  $F^+$  contains  $F \cup \{A \rightarrow C, B \rightarrow A, C \rightarrow B\}$ , so...
  - $F_{AB} \supseteq \{A \rightarrow B, B \rightarrow A\}$ ;  $F_{BC} \supseteq \{B \rightarrow C, C \rightarrow B\}$
  - So,  $(F_{AB} \cup F_{BC})^+ \supseteq \{C \rightarrow A\}$

# Third Normal Form (3NF)

- Reln R with FDs  $F$  is in 3NF if, for all  $X \rightarrow A$  in  $F^+$   
 $A \in X$  (called a *trivial* FD), or  
 $X$  is a superkey of R, or  
 $A$  is part of some *candidate* key (not superkey!) for R.  
(sometimes stated as “A is *prime*”)
- *Minimality* of a candidate key is crucial in third condition above!
- If R is in BCNF, obviously in 3NF.
- If R is in 3NF, some redundancy is possible. It is a compromise, used when BCNF not achievable (e.g., no “good” decomp, or performance considerations).
  - *Lossless-join, dependency-preserving decomposition of R into a collection of 3NF relations always possible.*



# Minimal Cover for a Set of FDs

- Minimal cover G for a set of FDs F:
  - Closure of F = closure of G.
  - Right hand side of each FD in G is a single attribute.
  - If we modify G by deleting an FD or by deleting attributes from an FD in G, the closure changes.
- Intuitively, every FD in G is needed, and "*as small as possible*" in order to get the same closure as F.
- e.g.,  $A \rightarrow B$ ,  $ABCD \rightarrow E$ ,  $EF \rightarrow GH$ ,  $ACDF \rightarrow EG$  has the following minimal cover:
  - $A \rightarrow B$ ,  $ACD \rightarrow E$ ,  $EF \rightarrow G$  and  $EF \rightarrow H$
- M.C. implies Lossless-Join, Dep. Pres. Decomp!!!
  - (in book)

# Concurrency Control

- Transaction: basic unit of operation
  - made up of reads and writes
- Goal: ACID Transactions
- A & D are provided by Crash Recovery
- C & I are provided by Concurrency Control
- Bottom line: reads and writes for various transactions **MUST** be ordered such that the final state of the database is the same as *some serial ordering of the transactions*

# Approaches to Concurrency Control

- 2PL - all objects have Shared and eXclusive locks
  - once one lock is released, no more locks may be acquired
  - Strict 2PL: don't release locks until commit time
- Locking issues
  - must either prevent or detect deadlock
  - may want multiple granularity locks (table, page, record) using IS, IX, SIX, S, X locks (check compatibility matrix!)
  - locking in B-trees usually not 2PL
  - phantom problem: locking all records of a given criteria (e.g., age > 20)

# Crash Recovery

- ACID - need way to ensure A & D
- We studied approach of Aries system
- Buffer management Steal, no Force
- Every Write to a page is first logged in WAL
  - log record is in stable storage before data page on disk
  - log record has Xact#, before value, after value
- Checkpoints record which pages dirty, which XActs running

---

# Transaction Commit

- write Commit record to log
  - flush log tail to stable storage
  - remove Xact from Xact table
  - write End record to log
-

---

# Transaction Abort

- write Abort record to log
  - go back through log, undoing each write (and add CLR to log)
  - when done, write End record to log
-

# Crash Recovery - 3 phases

- Analysis: starting from checkpoint, go forward in the log to see:
  - what pages were dirty
  - what transactions were active at time of crash
- Redo: start from oldest log record (oldest recLSN) that wrote to a dirty page, and redo all writes to dirty pages.
- Undo: start at the end of the log (time of crash), work backward undoing all writes made by transactions that were active at time of crash