

Implementation of Relational Operations

courtesy of Joe Hellerstein, Mike Franklin, and etc for some slides.

Jianlin Feng
School of Software
SUN YAT-SEN UNIVERSITY

Introduction

- Next topic: QUERY PROCESSING
- Some database operations are EXPENSIVE
- Huge performance gained by being “smart”
 - We’ll see 1,000,000x over naïve approach
- Main weapons are:
 - clever implementation techniques for operators
 - exploiting relational algebra “equivalences”
 - using statistics and cost models to choose


Simple SQL Refresher

- `SELECT <list-of-fields>`
 `FROM <list-of-tables>`
 `WHERE <condition>`

```
SELECT S.name, E.cid  
      FROM Students S, Enrolled E  
      WHERE S.sid=E.sid AND E.grade='A'
```

A Really Bad Query Optimizer

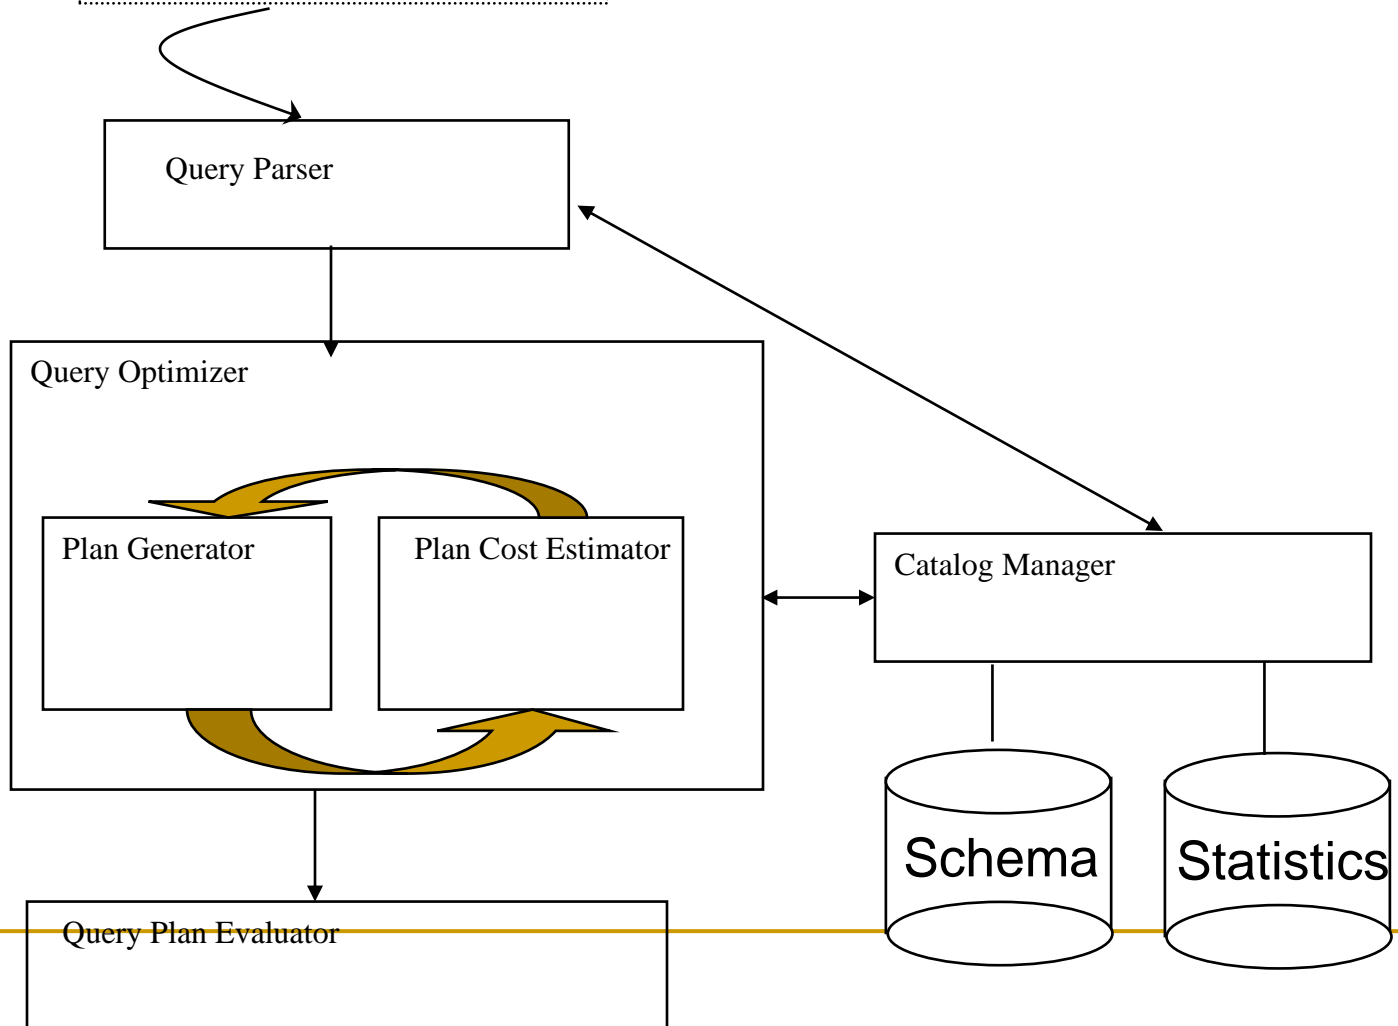
- For each Select-From-Where query block
 - Create a plan that:
 - Forms the cross product of the FROM clause
 - Applies the WHERE clause
- (Then, as needed:
 - Apply the GROUP BY clause
 - Apply the HAVING clause
 - Apply any projections and output expressions
 - Apply duplicate elimination and/or ORDER BY)



Cost-based Query Sub-System

Queries

```
Select *  
From Blah B  
Where B.blah = blah
```



The Query Optimization Game

- Goal is to pick a “good” plan
 - Good = low expected cost, under *cost model*
 - Degrees of freedom:
 - access methods
 - physical operators
 - operator orders
- Roadmap for this topic:
 - *First:* implementing individual operators
 - *Then:* optimizing multiple operators

Relational Operations

- We will consider how to implement:
 - Selection (σ) Select a subset of rows.
 - Projection (π) Remove unwanted columns.
 - Join (\bowtie) Combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 and in reln. 2.
- Q: What about Intersection?

Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

■ Sailors:

- ❑ Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
- ❑ $[S]=500$, $p_S=80$.

■ Reserves:

- ❑ Each tuple is 40 bytes, 100 tuples per page, 1000 pages.
- ❑ $[R]=1000$, $p_R=100$.

Simple Selections

```
SELECT *  
FROM Reserves R  
WHERE R.rname < 'C%'
```

$$\sigma_{R.attr op value}(R)$$

- How best to perform? Depends on:
 - what indexes are available
 - expected size of result
- **Size of result** approximated as
$$(size\ of\ R) * selectivity$$
 - *selectivity* estimated via statistics – we will discuss shortly.

Our options ...

- If no appropriate index exists:

Must scan the whole relation

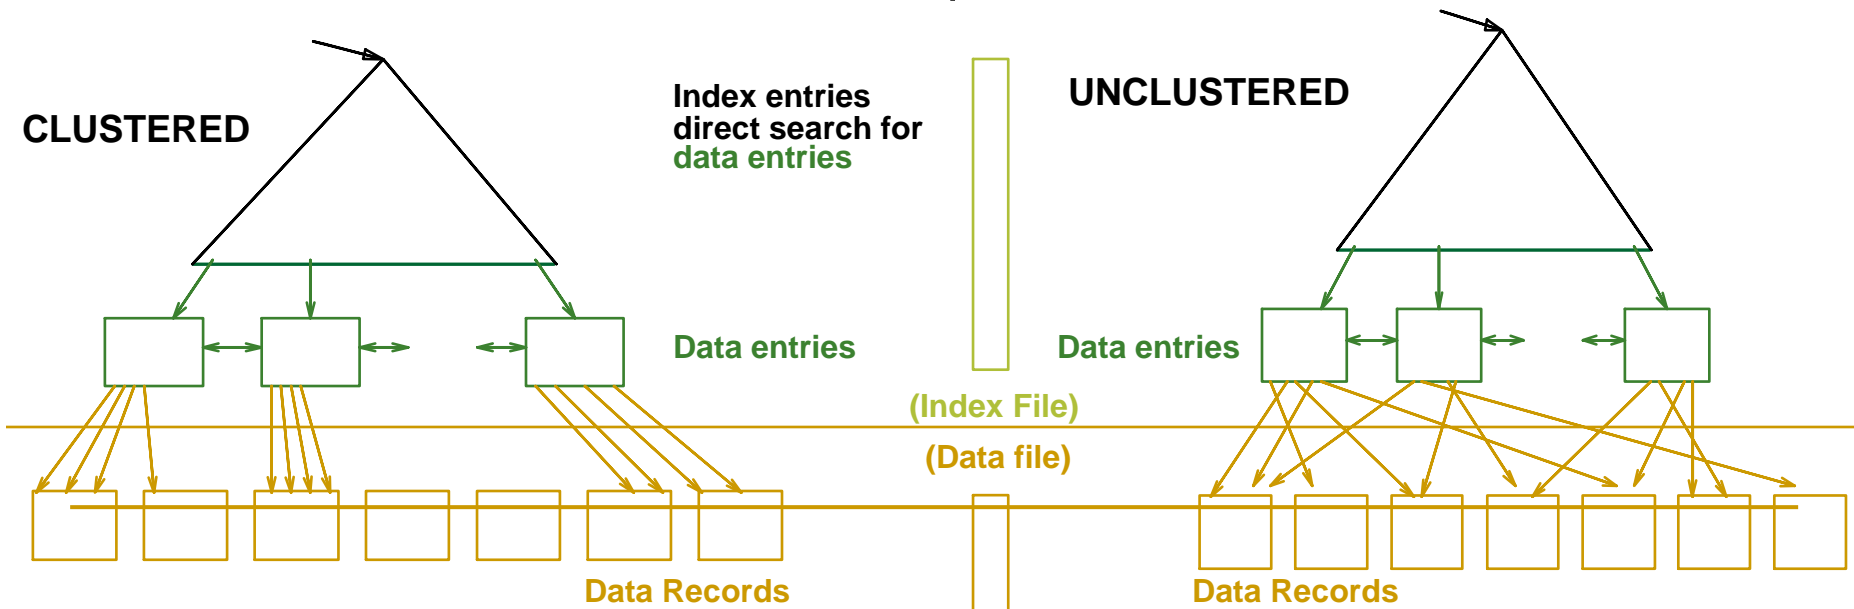
$\text{cost} = [R]$. For “reserves” = 1000 I/Os.

Our options ...

- With index on selection attribute:
 1. Use index to find qualifying data entries
 2. Retrieve corresponding data records

Total cost = cost of step 1 + cost of step 2

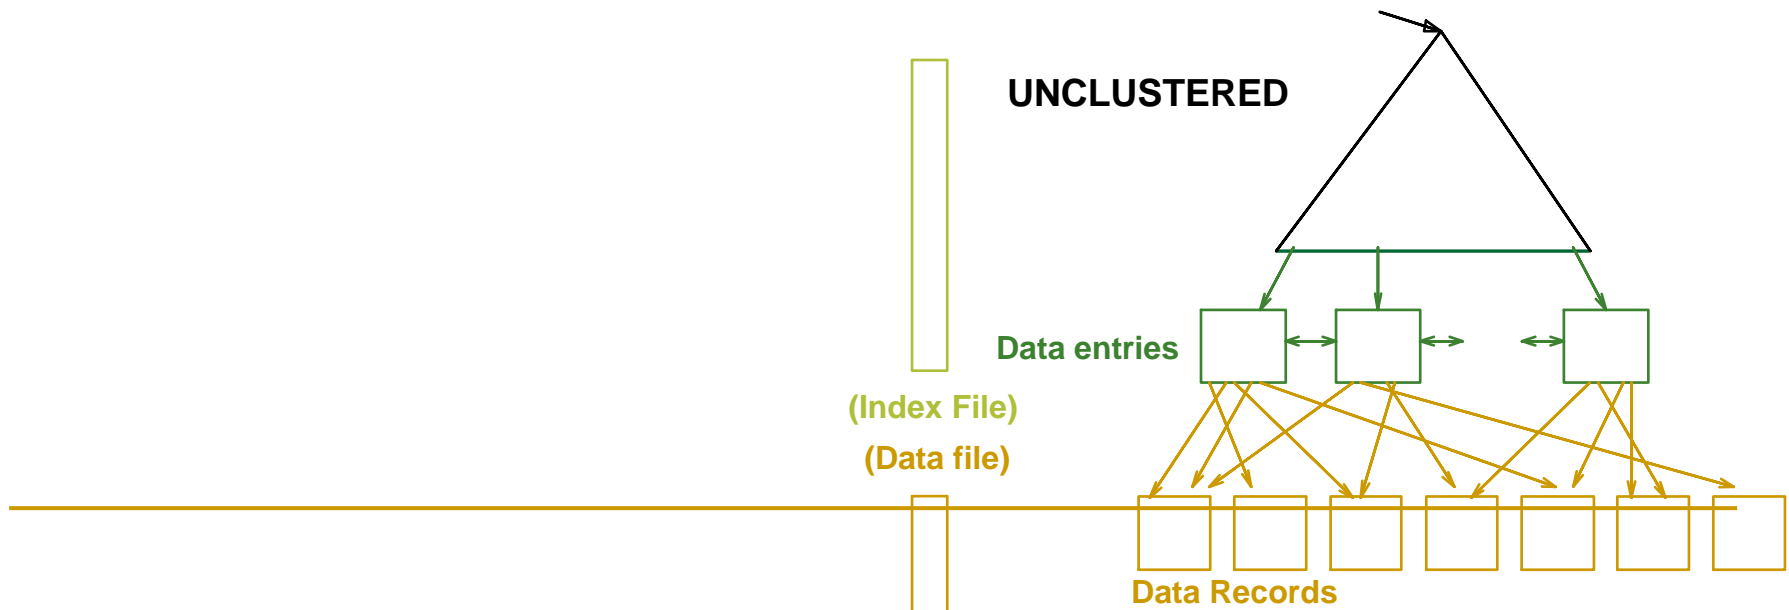
- For “reserves”, if selectivity = 10% (100 pages, 10000 tuples):
 - If *clustered* index, cost is a little over 100 I/Os;
 - If *unclustered*, could be up to 10000 I/Os! ... unless ...



Refinement for unclustered indexes

1. Find qualifying data entries.
2. Sort the **rids** of the data records to be retrieved.
3. Fetch rids in order.

Each data page is looked at just once (though # of such pages likely to be higher than with clustering).



General Selection Conditions

(day<8/9/94 AND rname='Paul') OR bid=5 OR sid=3

- First, convert to conjunctive normal form (CNF):
 - *(day<8/9/94 OR bid=5 OR sid=3) AND (rname='Paul' OR bid=5 OR sid=3)*
- We only discuss the case with no ORs
- Terminology:
 - A **B-tree** index matches terms that involve only attributes in a *prefix* of the search key. e.g.:
 - Index on *<a, b, c>* matches *a=5 AND b= 3*, but not *b=3*.

2 Approaches to General Selections

Approach I:

1. Find the *cheapest access path*
 2. retrieve tuples using it
 3. Apply any remaining terms that don't **match** the index
-
- ❑ *Cheapest access path*: An index or file scan that we estimate will require the fewest page I/Os.

Cheapest Access Path - Example

query: *day < 8/9/94 AND bid=5 AND sid=3*

some options:

B+tree index on day; check bid=5 and sid=3 afterward.

hash index on <bid, sid>; check day<8/9/94 afterward.

- *How about a B+tree on <rname, day>?*
- *How about a B+tree on <day, rname>?*
- *How about a Hash index on <day, rname>?*

2 Approaches to General Selections (Contd.)

Approach II: use 2 or more matching indexes.

1. From each index, get set of rids
2. Compute intersection of rid sets
3. Retrieve records for rids in intersection
4. Apply any remaining terms

EXAMPLE: *day<8/9/94 AND bid=5 AND sid=3*

Suppose we have an index on *day*, and another index on *sid*.

- ❑ Get rids of records satisfying *day<8/9/94*.
 - ❑ Also get rids of records satisfying *sid=3*.
 - ❑ Find intersection, then retrieve records, then check *bid=5*.
-

Projection

```
SELECT  DISTINCT
        R.sid, R.bid
FROM    Reserves R
```

- Issue is removing **duplicates**.
pass 1: $\lceil \frac{250}{20} \rceil = 13$ 12个 20
1个 10
- Use sorting!!
pass 2: 剩下 7个 buffer
可以满足 13路
归并排序
- 1. Scan R, extract only the needed attributes
- 2. Sort the resulting set 外排序
- 3. Remove adjacent duplicates

Cost:

Ramakrishnan/Gehrke writes to temp table at each step!

Reserves with size ratio $0.25 = 250$ pages.

With 20 buffer pages can sort in 2 passes, so:

$$1000 + 250 + 2 * 2 * 250 + 250 = 2500 \text{ I/Os}$$

扫描 取出 2趟外排序 消重

Projection -- improved

- Avoid the temp files, work on the fly:
 - Modify Pass 0 of sort to eliminate unwanted fields.
 - Modify Passes 1+ to eliminate duplicates.

Cost:

Reserves with size ratio 0.25 = 250 pages.

With 20 buffer pages can sort in 2 passes, so:

1. Read 1000 pages
2. Write 250 (in runs of 40 pages each) = 7 runs
3. Read and merge runs (20 buffers, so 1 merge pass!)

Total cost = 1000 + 250 + 250 = 1500.

Other Projection Tricks

If an index search key contains all wanted attributes:

- Do *index-only* scan
 - Apply projection techniques to data entries (*much smaller!*)

If a B+Tree index search key *prefix* has all wanted attributes:

- Do *in-order index-only* scan
 - Just retrieve the data entries in order;
 - Discarding unwanted fields;
 - Compare adjacent tuples on the fly *to check for duplicates*.

Joins

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- Joins are very common.
- $R \bowtie S$ is large; so, $R \bowtie S$ followed by a selection is inefficient.
- Many approaches to reduce join cost.
- Join techniques we will cover today:
 1. Nested-loops join
 2. Index-nested loops join
 3. Sort-merge join

Simple Nested Loops Join

$R \bowtie S$: foreach tuple r in R do

直接读元组

foreach tuple s in S do

if $r_i == s_j$ then add $\langle r, s \rangle$ to result

$$\text{Cost} = (p_R * [R]) * [S] + [R] = 100 * 1000 * 500 + 1000 \text{ IOs}$$

□ At 10ms/IO, Total time: ??? 比较 S读入 R读入

只需要3个缓冲区

- What if smaller relation (S) was “outer”?
- What assumptions are being made here?
- What is cost if one relation can fit entirely in memory?

Page-Oriented Nested Loops Join

$R \bowtie S$:

读页面

```
foreach page  $b_R$  in R do
  foreach page  $b_S$  in S do
    foreach tuple  $r$  in  $b_R$  do
      foreach tuple  $s$  in  $b_S$  do
        if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```


$$\text{Cost} = [R] * [S] + [R] = 1000 * 500 + 1000$$

outer X inner + outer

- If smaller relation (S) is outer, cost = $500 * 1000 + 500$
- Much better than naïve per-tuple approach!

Block Nested Loops Join

- Page-oriented NL doesn't exploit extra buffers :(
- Idea to use memory efficiently:



Cost: Scan outer + (#outer blocks * scan inner)

$\#outer\ blocks = \left\lceil \frac{\#of\ pages\ of\ outer}{blocksize} \right\rceil$

Examples of Block Nested Loops Join

块嵌套

2个buffer进行输入输出

- Say we have $B = 100+2$ memory buffers
- Join cost = [outer] + (#outer blocks * [inner])
#outer blocks = [outer] / 100
- With R as outer ($[R] = 1000$):
 - Scanning R costs 1000 IO's (*done in 10 blocks*)
 - Per block of R, we scan S; costs 10×500 I/Os
 - **Total = $1000 + 10 \times 500$.**
- With S as outer ($[S] = 500$):
 - Scanning S costs 500 IO's (*done in 5 blocks*)
 - Per block of S, we scan R; costs 5×1000 IO's
 - ~~Total = $500 + 5 \times 1000$.~~


Index Nested Loops Join 索引嵌套

$R \bowtie S$: foreach tuple r in R do

foreach tuple s in S where $r_i == s_j$ do

add $\langle r, s \rangle$ to result

R : r_i



Index Nested Loops Join

```
R ⋈ S: foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add <r, s> to result
```

Cost = $[R] + ([R] * p_R) * \text{cost to find matching S tuples}$



- If index uses **Alt. 1**, cost = cost to traverse tree from root to leaf.
- For **Alt. 2 or 3**:
 1. Cost to lookup RID(s); typically 2-4 IO's for B+Tree.
 2. Cost to retrieve records from RID(s); depends on clustering.
 - ❑ **Clustered index**: 1 I/O per page of matching S tuples.
 - ❑ **Unclustered**: up to 1 I/O per matching S tuple.

Reminder: Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

■ Sailors:

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
- $[S]=500$, $p_S=80$.

■ Reserves:

- Each tuple is 40 bytes, 100 tuples per page, 1000 pages.
- $[R]=1000$, $p_R=100$.

Sort-Merge Join

外排序 归并

Example:

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

1. Sort R on join attr(s)
2. Sort S on join attr(s)
3. Scan sorted-R and sorted-S in tandem, to find matches

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

<u>sid</u>	<u>bid</u>	<u>day</u>	rname
28	103	12/4/96	guppy
28	103	11/3/96	yuppy
31	101	10/10/96	dustin
31	102	10/12/96	lubber
31	101	10/11/96	lubber
58	103	11/12/96	dustin

Cost of Sort-Merge Join

- Cost: $\text{Sort } R + \text{Sort } S + ([R] + [S])$
 - But in worst case, last term could be $[R] * [S]$ (*very unlikely!*)
 - Q: what is worst case?

Suppose $B = 35$ buffer pages:

- Both R and S can be sorted in 2 passes
- Total join cost = $4 * 1000 + 4 * 500 + (1000 + 500) = 7500$

$$\left\lceil \frac{1000}{35} \right\rceil < 35 \quad \left\lceil \frac{500}{35} \right\rceil < 35$$

都只需要 2 趟外排序

Suppose $B = 300$ buffer pages:

- Again, both R and S sorted in 2 passes
- Total join cost = 7500

Block-Nested-Loop cost = 2500 ... 15,000

Refinement of Sort-Merge Join

- We can combine the merging phases in the *sorting* of R and S with the merging required for the join.
 - If $B > \sqrt{L}$, where L is the size of the **larger** relation,
 - using the sorting refinement (13.3.1) that produces runs of length $2B$ in Pass 0, #runs of each relation is $< B/2$.
- In “Merge” phase: Allocate 1 page per run of **each relation**, and ‘merge’ while checking the join condition
- **Cost:** read+write each relation in Pass 0 + read each relation in (only) merging pass (+ writing of result tuples).
- In example, cost goes down from 7500 to 4500 I/Os.

Hash-Join

口合X
口布

- **Partition** both relations on the join attributes using hash function **h**.
- R tuples in partition R_i will **only** match S tuples in partition S_i .


For $i = 1$ to #partitions {

Read in partition R_i
and hash it using
h2 (not **h**).


Scan partition S_i and
probe hash table
for matches.

}

Original
Relation




Disk



B main memory buffers

Partitions




1

2


B-1

Disk

Partitions
of R & S




Disk



B main memory buffers

Join Result



Disk

Memory Requirements of Hash-Join

- #partitions $k < B$, and $B-1 > \text{size of largest partition}$ to be held in memory. Assuming uniformly sized partitions, and maximizing k , we get:
 $k = B-1$, and $M/(B-1) < B-2$, i.e., B must be $> \sqrt{M}$
从用小的表
- Since we build an in-memory hash table to speed up the matching of tuples in the second phase, a little more memory is needed.
- If the hash function does not partition uniformly, one or more R partitions may not fit in memory.
 - Can apply hash-join technique **recursively** to do the join of this R -partition with corresponding S -partition.

Cost of Hash-Join

$$3 \times (M+N) = 3 \times (1000 + 500)$$

- In partitioning phase, read+write both relns; $2(M+N)$.
In matching phase, read both relns; $M+N$ I/Os.

$$\sqrt{500} = 23 \quad \left\lceil \frac{1000}{23} \right\rceil = 44 \quad \frac{44}{22} = 2$$

- In our running example, this is a total of 4500 I/Os.

$$2 \times 3 \times 1000 + 2 \times 2 \times 500 + 1000 + 500 = 9500 > 4500$$

- Sort-Merge Join vs. Hash Join:

- Given a minimum amount of memory (*what is this, for each?*) both have a cost of $3(M+N)$ I/Os. Hash Join superior if relation sizes differ greatly (e.g., if one reln fits in memory). Also, Hash Join shown to be highly parallelizable.
- ~~Sort-Merge less sensitive to data skew; result is sorted.~~

Set Operations

- Intersection and cross-product as special cases of join.
- Union (Distinct) and Except similar; we'll do union.
- Sorting based approach to union:
 - Sort both relations (on combination of all attributes).
 - Scan sorted relations and merge them.
 - *Alternative:* Merge runs from Pass 0 for *both* relations.
- Hash based approach to union: 用主键进行hash
 - Partition R and S using hash function h .
 - For each S-partition, build in-memory hash table (using h_2), scan corresponding R-partition and add tuples to table while discarding duplicates.

General Join Conditions

- Equalities over several attributes (e.g., *R.sid=S.sid* AND *R.rname=S.sname*):
 - For Index NL, build index on *<sid, sname>* (if S is inner); or use existing indexes on *sid* or *sname*.
 - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.
- Inequality conditions (e.g., *R.rname < S.sname*):
 - For Index NL, need (clustered!) B+ tree index.
 - Range probes on inner; # matches likely to be much higher than for equality joins.
 - Hash Join, Sort Merge Join not applicable!
 - Block NL quite likely to be the best join method here.

Aggregate Operations (AVG, MIN, etc.)

Example:

```
SELECT AVG(S.age)  
FROM Sailors S
```

- Without grouping:
 - In general, requires scanning the relation.
 - Given a tree index whose search key includes all attributes in the **SELECT** or **WHERE** clauses, can do index-only scan.

Aggregate Operations (continued)

■ With grouping:

- ❑ Sort on group-by attributes, then scan relation and compute aggregate for each group. (Better: combine sorting and aggregate computation.)
- ❑ Similar approach based on hashing on group-by attributes.
- ❑ Given a tree index whose search key includes all attributes in SELECT, WHERE and GROUP BY clauses, can do index-only scan;
 - if group-by attributes form prefix of search key, can retrieve data entries/tuples in group-by order.

Summary

- *Queries are composed of a few basic operators;*
 - The implementation of these operators can be carefully tuned (and it is important to do this!).
- Many alternative implementation techniques for each operator; no universally superior technique for most.
- Must consider alternatives for each operation in a query and choose best one based on statistics, etc.
- This is part of the broader task of Query Optimization, which we will cover next!