

Physical DB Design

courtesy of Joe Hellerstein for some slides

Jianlin Feng

School of Software

SUN YAT-SEN UNIVERSITY

Physical DB Design

- Query optimizer does what it can to use indices, clustering etc.
 - DataBase Administrator (DBA) is expected to set up physical design well.
 - Good DBAs understand query optimizers very well.
-

One Key Decision: Indexes

- Which tables
 - Which field(s) should be the search key?
 - Multiple indexes?
 - Clustering?
-

Index Selection

- One approach:
 - ❑ Consider most important queries in turn.
 - ❑ Consider best plan using the current indexes
 - ❑ See if better plan is possible with an additional index.
 - ❑ If so, create it.
- But consider impact on updates!
 - ❑ Indexes can make queries go faster, updates slower.
 - ❑ Require disk space, too.

Issues to Consider in Index Selection

- Attributes mentioned in a `WHERE` clause are candidates for index search keys.
 - Range conditions are sensitive to clustering
 - Exact match conditions don't require clustering
 - Or do they???? :-)
- Choose indexes that benefit many queries
- NOTE: only one index can be clustered per relation!
 - So choose it wisely!

Example 1

```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE E.dno=D.dno AND D.dname='Toy'
```

- B+ tree index on *D.dname* supports 'Toy' selection.
 - Given this, index on *D.dno* is not needed.
- B+ tree on *E.dno* allows us to get matching (inner) Emp tuples for each selected (outer) Dept tuple.
- What if WHERE included: `` ... AND E.age=25'' ?
 - Could retrieve Emp tuples using index on *E.age*, then join with Dept tuples satisfying *dname* selection.
 - Comparable to strategy that used *E.dno* index.
 - So, if *E.age* index is already created, this query provides much less motivation for adding an *E.dno* index.

E.dno = 007 , *D.dno* 索引 | B+ tree

Example 2

聚集索引 利于范围查询

```
SELECT E.ename, D.mgr
FROM Emp E, Dept D
WHERE E.sal BETWEEN 10000 AND 20000
AND E.hobby='Stamps' AND E.dno=D.dno
```

范围查询

- All selections are on Emp so it should be the outer relation in any Index NL join. 嵌套循环，条件选择后为外层
 - Suggests that we build a B+ tree index on *D.dno*.
- What index should we build on Emp?
 - B+ tree on *E.sal* could be used, OR an index on *E.hobby* could be used.
 - Only one of these is needed, and which is better depends upon the selectivity of the conditions.
 - As a rule of thumb, equality selections more selective than range selections. 点查询优于范围查询，选择 E.hobby 建 B+ tree
- *Have to understand optimizers to get this right!*

Examples of Clustering

索引
选择范围越小收益越大

```
SELECT E.dno
FROM Emp E
WHERE E.age>40
```

单边范围查询

■ B+ tree index on E.age can be used to get qualifying tuples.

- How selective is the condition?
- Is the index clustered? ✓

```
SELECT E.dno, COUNT (*)
FROM Emp E
WHERE E.age>10
GROUP BY E.dno
```

■ Consider the GROUP BY query.

- If many tuples have E.age > 10, using E.age index and sorting the retrieved tuples may be costly. 选择范围过大
- Clustered E.dno index may be better!

■ Equality queries and duplicates:

- Clustering on E.hobby helps!

```
SELECT E.dno
FROM Emp E
WHERE E.hobby=Stamps
```